

CSC501

Operating Systems Principles

Memory Management

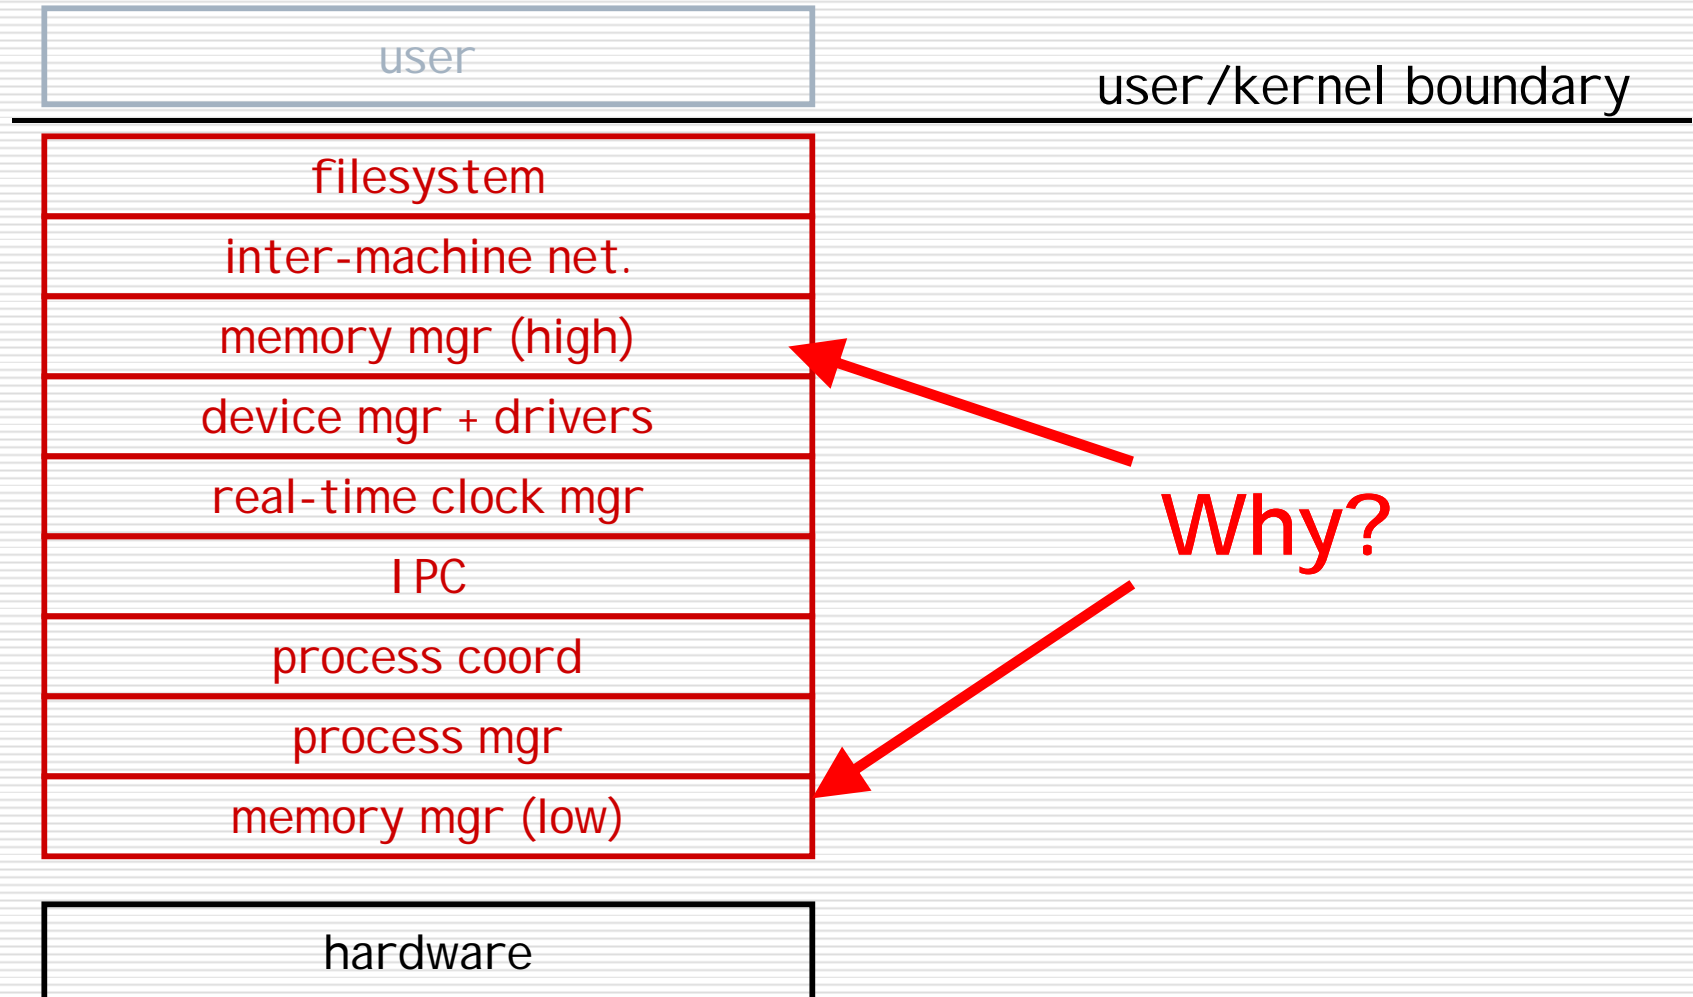
Previous Lectures

q Linker & Loader

q Today

Q Memory Management

Layer



Memory Management

- q Low-level memory management
 - Q Manages memory within kernel address space
 - Q Allocates address spaces for processes
 - Q Treats memory as a single, exhaustible resource
 - Q In layered-OS design, it is positioned in the hierarchy below process manager
- q High-level memory management
 - Q Manages pages within address space
 - Q Divides memory into abstract resources
 - Q Positioned in the hierarchy above device manager

Low-level Memory Management

q Conceptual Uses

Q Allocation of process stack

- ✓ *getstk, freestk*

- ✓ Used by OS

Q Allocation of heap storage

- ✓ *getmem, freemem*

- ✓ Used by applications and OS

Possible Allocation Strategies

- q Single free list
 - Q First-fit: scan free list and allocate first hole that is large enough
 - Q Next-fit: start search from end of last allocation
 - Q Best-fit: find smallest hole that is adequate
 - Q Worst fit: find largest hole
 - q Multiple lists
 - Q By exact size (static / dynamic)
 - q Hierarchy
 - Q Binary size allocation
 - q FIFO cache with above methods
-

Example Implementation in Xinu


- q Single free list
 - Q Ordered by increasing address
 - Q Singly-linked
 - Q Initialized at system startup to all free memory
- q Allocation policy
 - Q Heap: first-fit
 - Q Stack: last-fit
 - Q Minimizes fragmentation

Result of Xinu Allocation Policy



- q Stack allocation from highest free memory
- q Heap allocation from lowest free memory

Illustration of Xinu Free List



- q List in order of address
- q Related Files: *h/mem.h sys/getmem.c sys/freemem.c*

High-level Memory Management

- q Memory is cheap today, and getting cheaper
 - Q But applications are demanding more and more memory, there is never enough!
 - q Let's take a look at the history
-

In the Beginning (prehistory)...


- q Batch systems
 - Q One program loaded in physical memory
 - Q Runs to completion
 - q If job larger than physical memory, use **overlays**
 - Q Identify sections of program that
 - ✓ Can run to a result
 - ✓ Can fit into the available memory
 - Q Add statement after result to load a new section
 - Q Like passes in a compiler
-

In the Beginning (multi-programming) ...

- q Multiple processes in physical memory at the same time
 - Q Allows fast switching to a ready process
 - Q Divide physical memory into multiple pieces – **partitioning**
 - q Partition requirements
 - Q Protection – keep processes from smashing each other
 - Q Fast execution – memory accesses can't be slowed by protection mechanisms
 - Q Fast context switch – can't take forever to setup mapping of addresses
-

Physical Memory

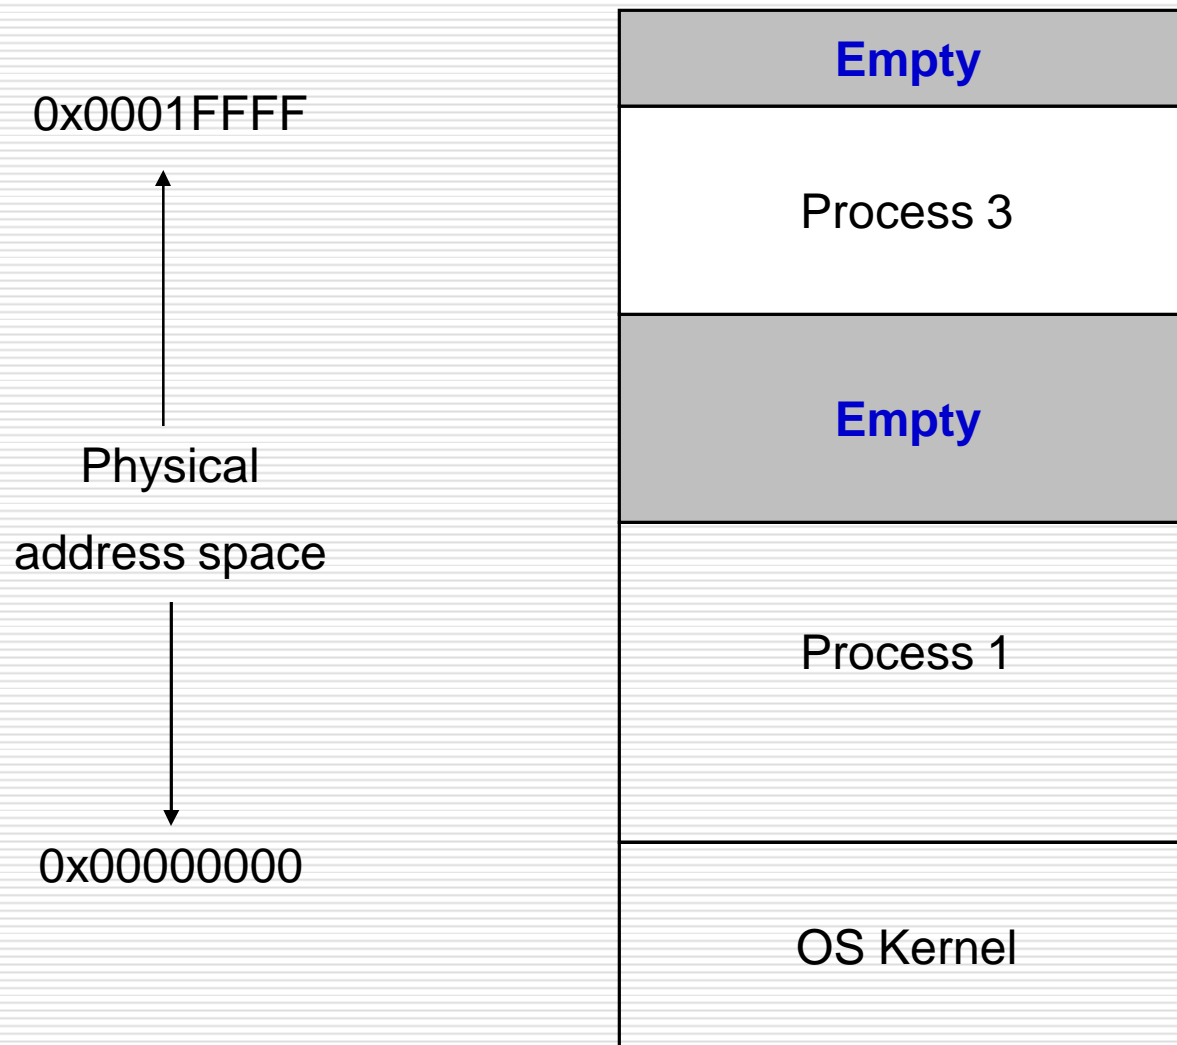
0x0000FFFF
↑
Physical
address space
↓
0x00000000



Loading a Process

- q Relocate all addresses relative to start of partition
 - Q See Linking and Loading
 - q Memory protection assigned by OS
 - Q Block-by-block to physical memory
 - q Once process starts
 - Q Partition cannot be moved in memory
 - Q Why?
-

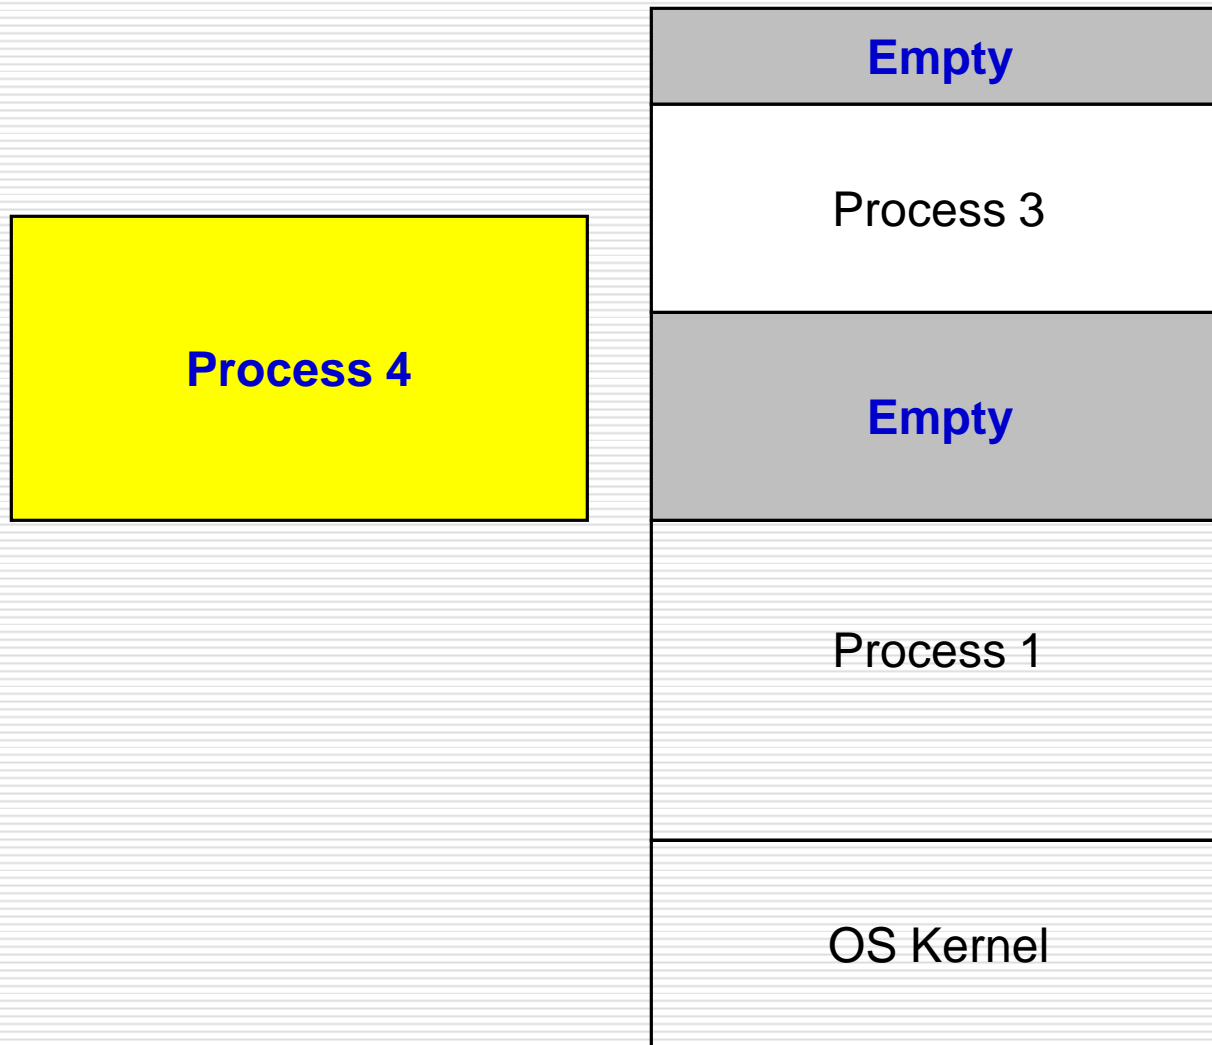
Physical Memory



Problem

- q What happens when Process 4 comes along and requires space larger than the largest empty partition?
 - ✓ Wait
 - ✓ Complex resource allocation problem for OS
 - ✓ Potential starvation
-

Physical Memory

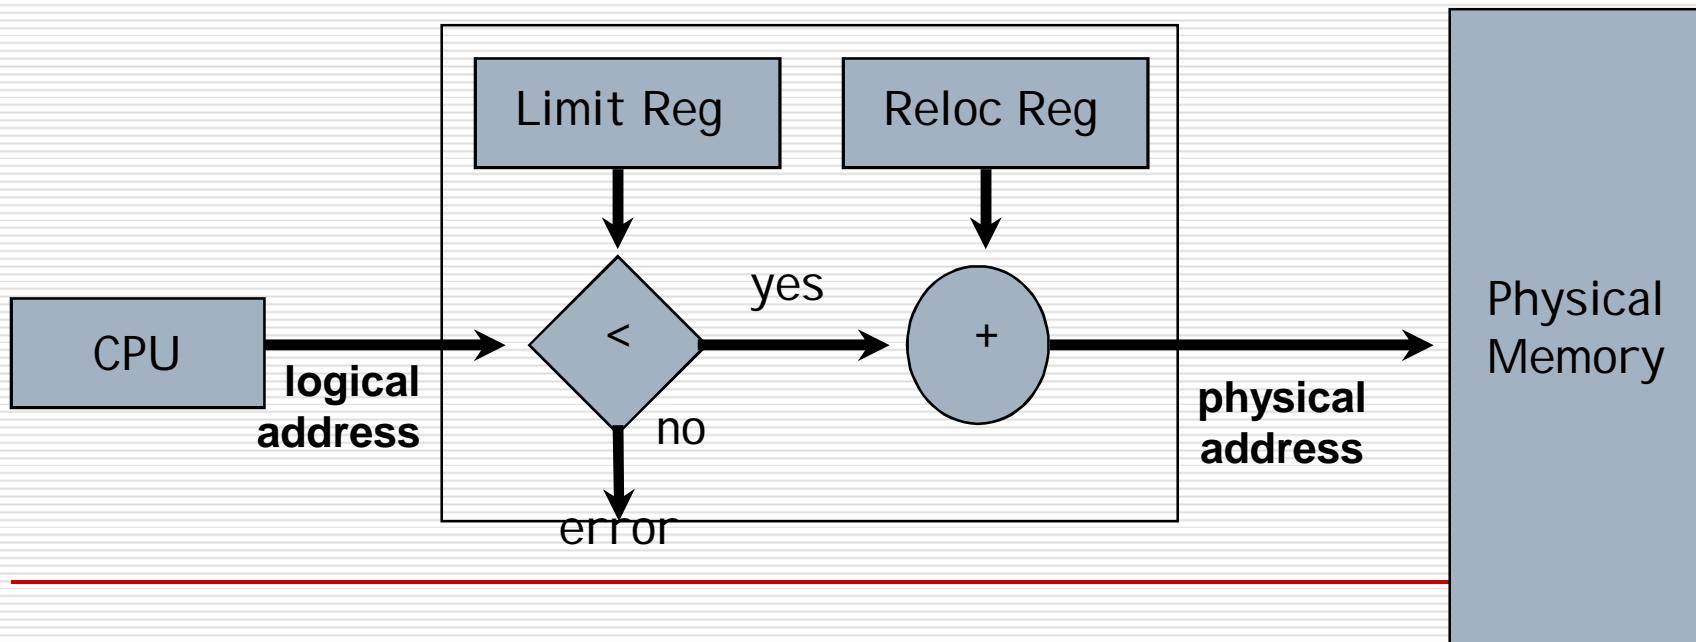


Solution

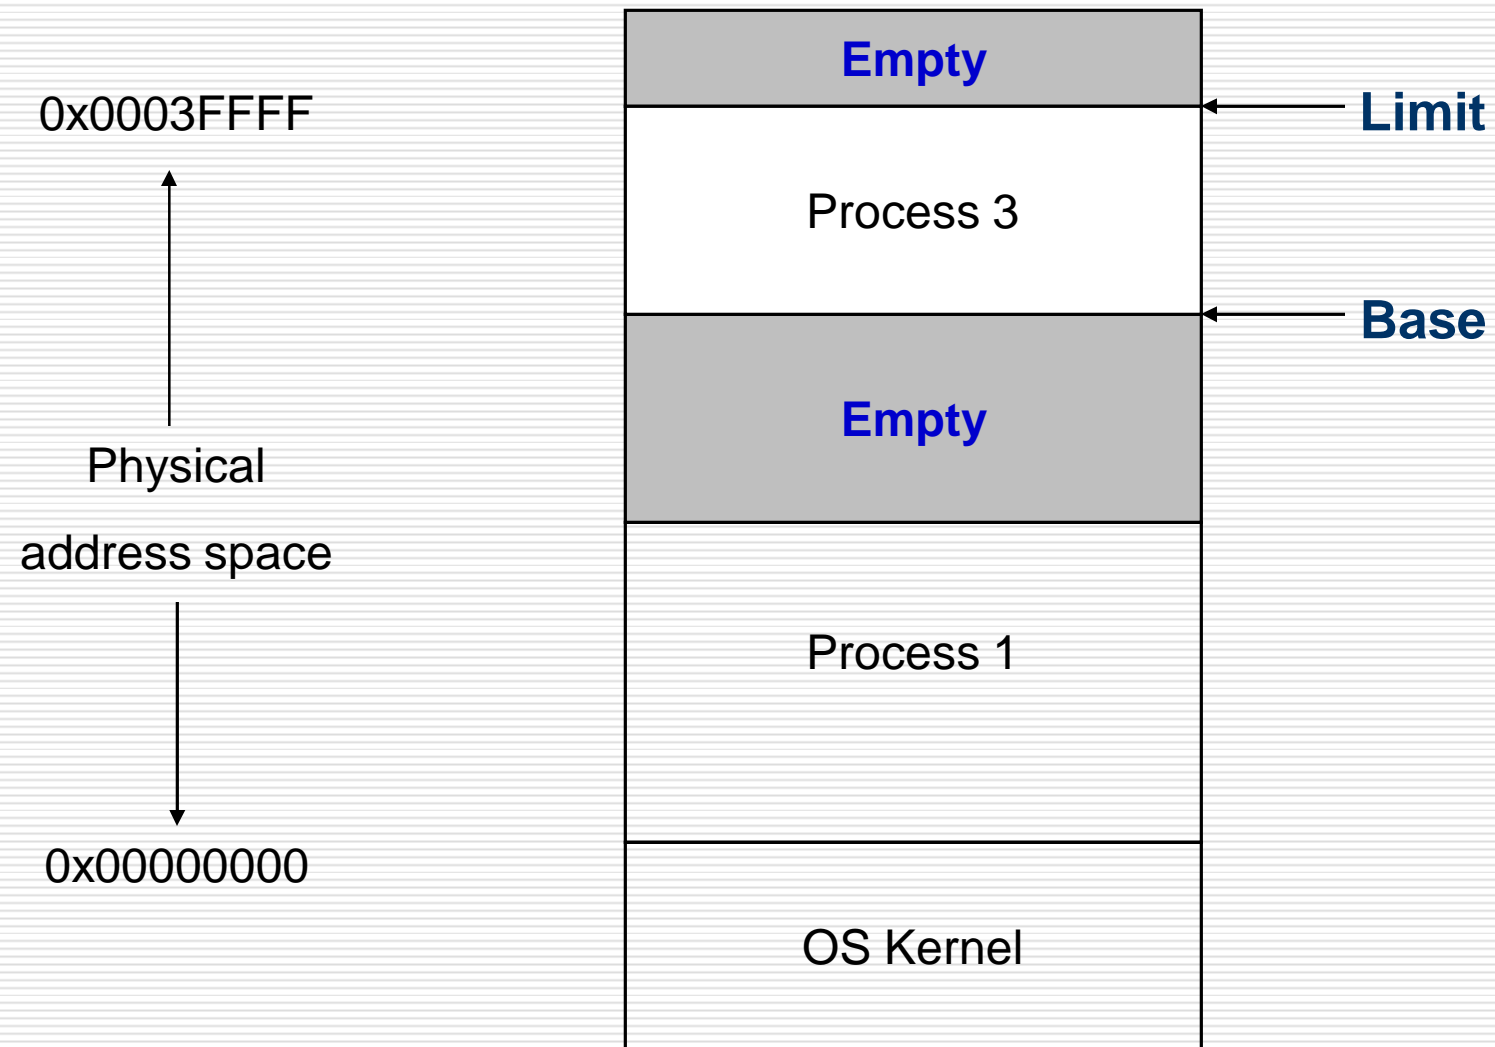
- q **Logical Address**: an address used by the program that is translated by computer into a **physical address** each time it is used
 - q When the program utters 0x00105C, ...
 - Q The machine accesses 0x01605C instead
-

Simplest Implementation

- q **Base** and **Limit** registers
 - Q *Base* added to all addresses
 - Q *Limit* checked on all memory references
- q Loaded by OS at each context switch



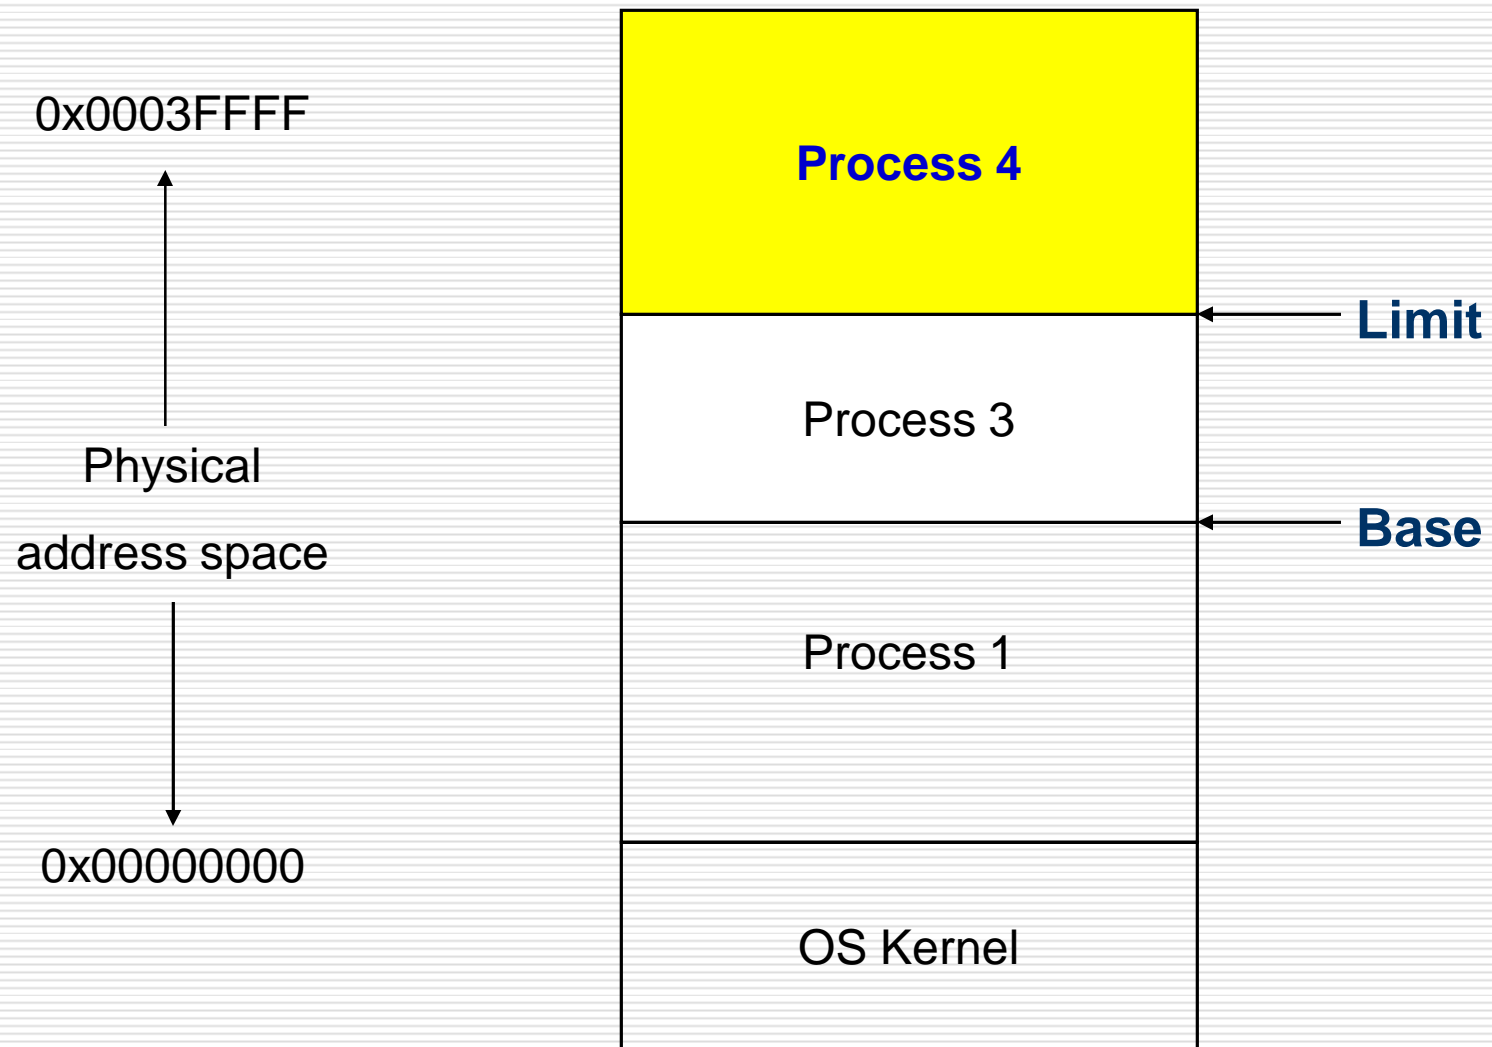
Physical Memory



Advantages

- q No relocation of program addresses at load time
 - Q All addresses relative to zero
 - q Built-in protection provided by *Limit*
 - Q No physical protection per page or block
 - q Fast execution
 - Q Addition and limit check at hardware speeds within each instruction
 - q Fast context switch
 - Q Need only change base and limit registers
 - q Partition can be suspended and moved at any time
 - Q Process is unaware of change
 - Q Expensive for large processes!
-

Physical Memory



Challenge – Memory Allocation

- q Fixed partitions
 - q Variable partitions
-

Partitioning Strategies – Fixed

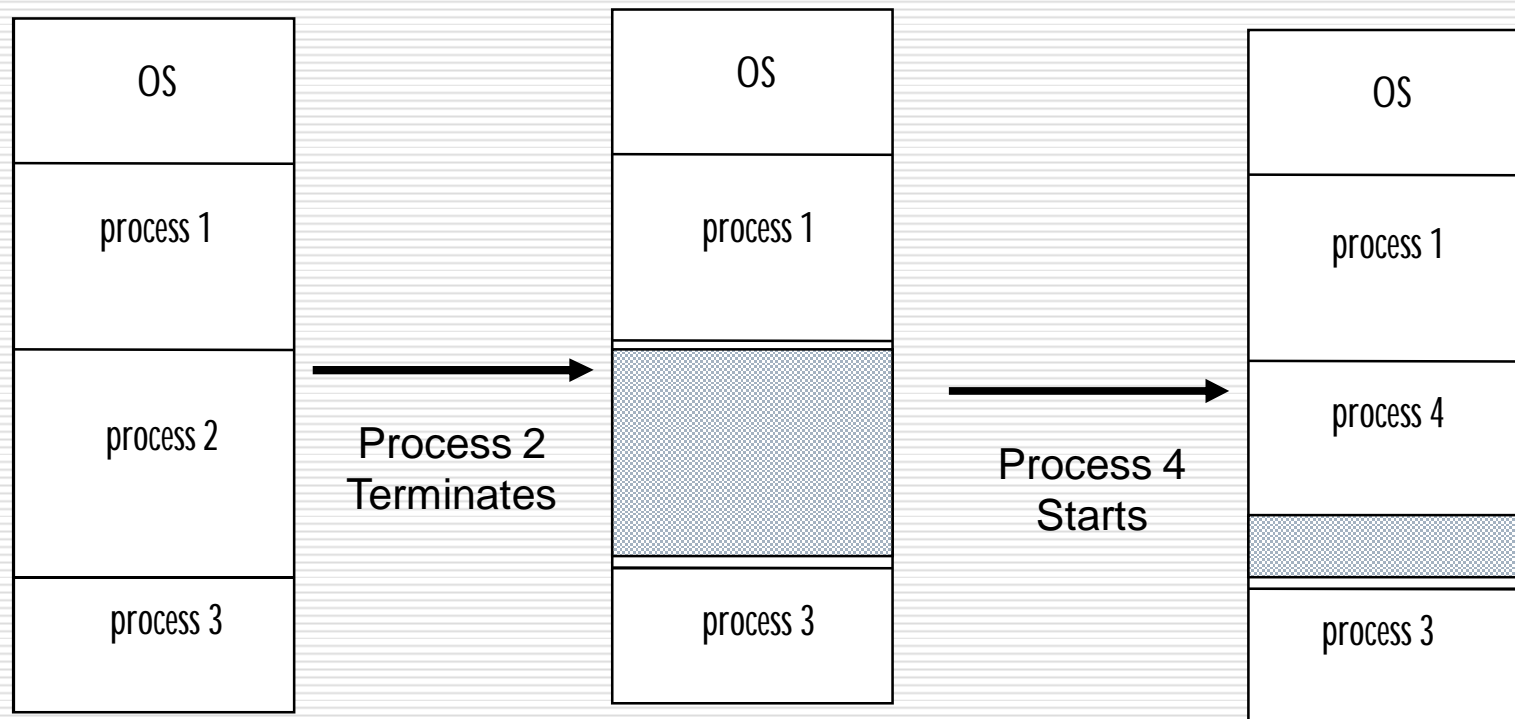
- q Fixed Partitions – divide memory into equal sized pieces (except for OS)
 - Q Degree of multiprogramming = number of partitions
 - Q Simple policy to implement
 - ✓ All processes must fit into partition space
 - ✓ Find any free partition and load the process
- q Problem – **Internal Fragmentation**
 - Q Unused memory in a partition that is not available to other processes

Question:
What is the “right”
partition size?

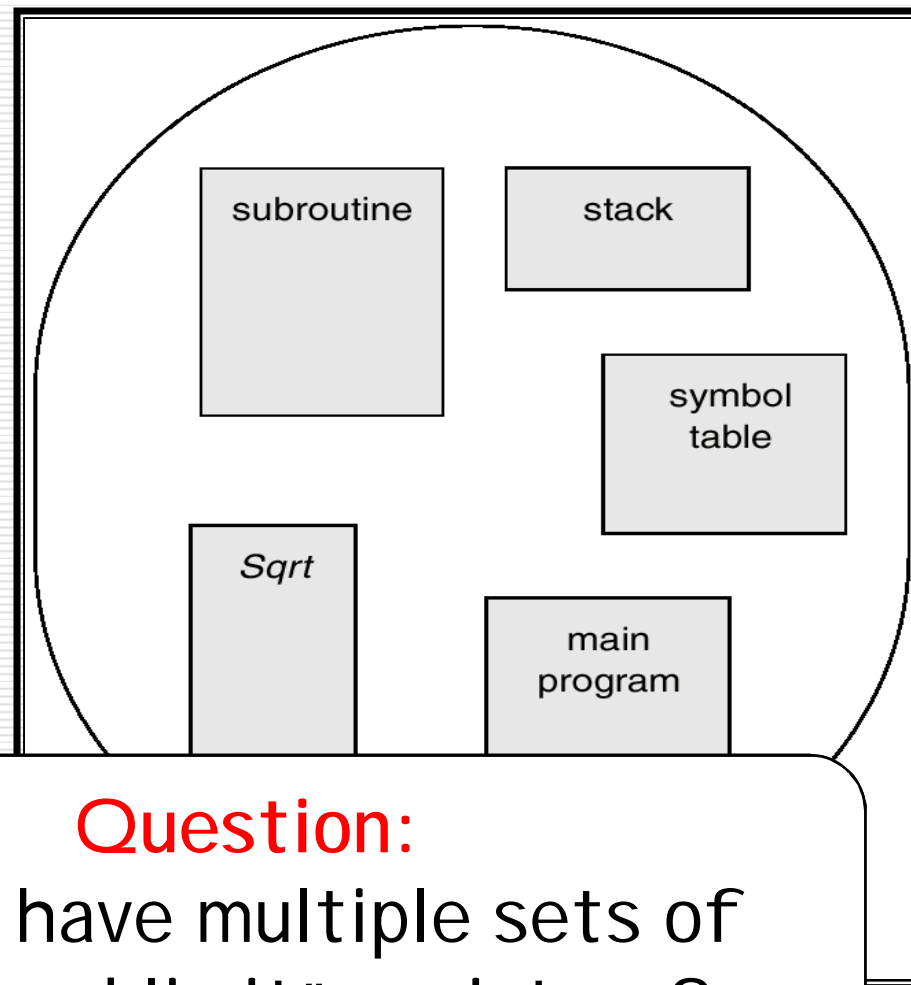
Partitioning Strategies – Variable

- q Memory is dynamically divided into partitions based on process needs
 - Q More complex management problem
 - ✓ Need data structures to do tracking of free and used memory
 - ✓ New process is allocated memory from hole large enough to fit it
 - q Problem – **External Fragmentation**
 - Q Unused memory between partitions that is not too small to be used by any processes
-

Partitioning Strategies – Variable




User's View of a Program




Question:

Can we have multiple sets of
"base and limit" registers?

Logical View of Segmentation



user space

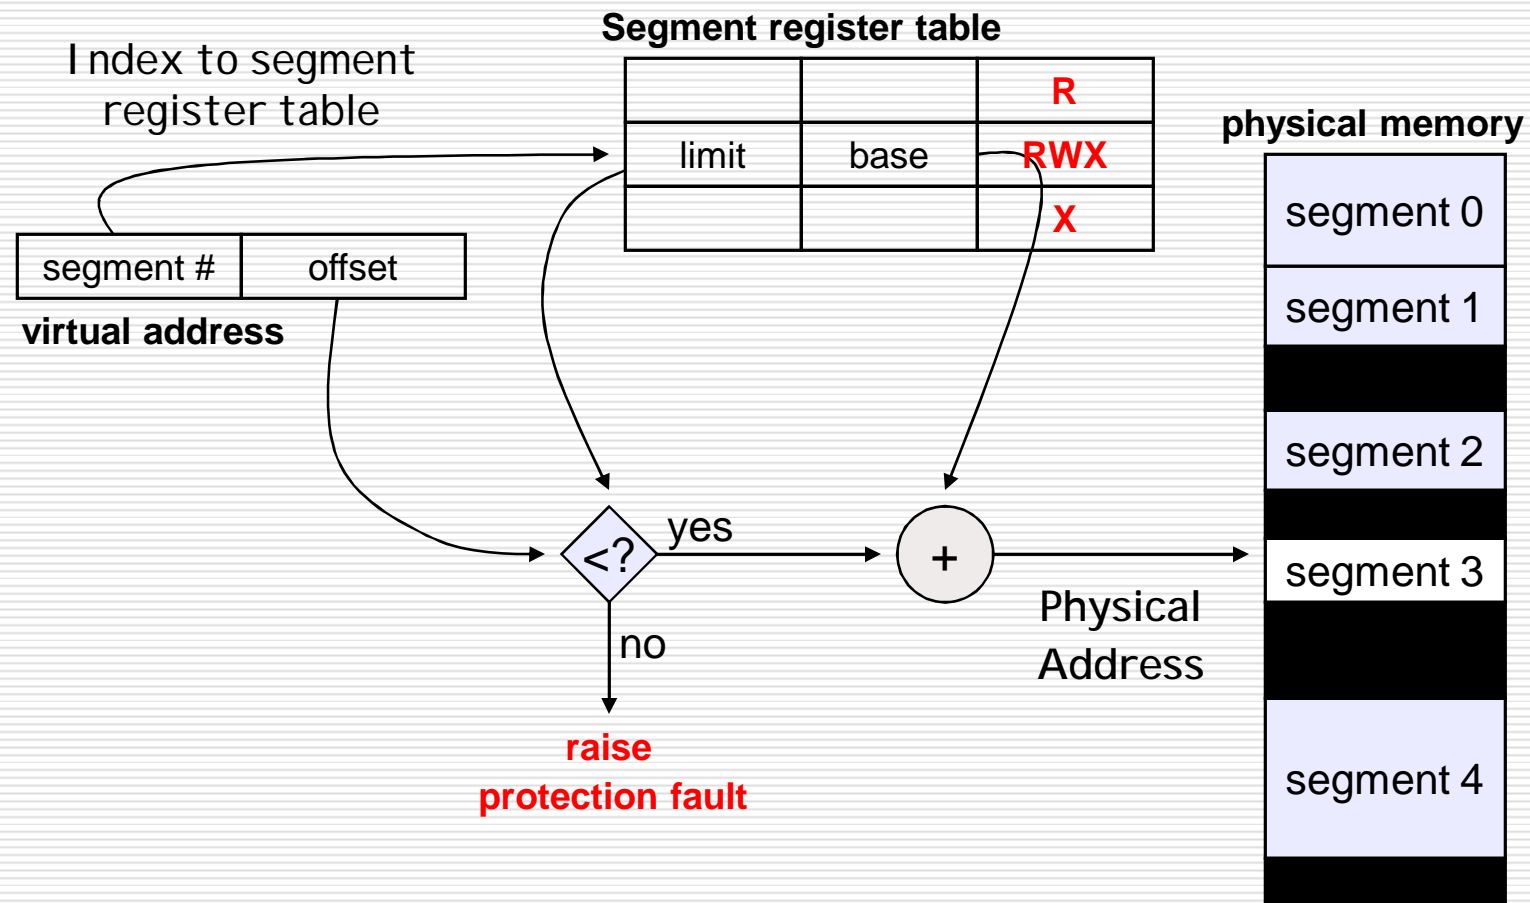


physical memory space

Segmentation

- q Logical address consists of a pair:
 <segment-number, offset>
 - q Segment table where each entry has:
 - Q Base: contains the starting physical address where the segments reside in memory.
 - Q Limit: specifies the length of the segment.
-

Segment Lookup



Segmentation

- q Common in early minicomputers
 - Q Small amount of additional hardware – 4 or 8 segments
 - Q Used effectively in Unix
- q Good idea that has persisted and supported in current hardware and OSs
 - Q X86 supports segments
 - Q Linux supports segments

Question:

Do we still have external fragmentation problem? If yes, can we further improve it?

Next Lecture

q Paging

Lab3 Out !